

# EECS6895 Adv. Bigdata Analytics HW 3 - Sun-Yi Lin(sl3833)

## Amazon EC2 + PyCUDA

Since I don't have a device with nVidia Graphic Card, so I choose to use Amazon AWS's Elastic Compute Cloud service plus PyCUDA to use GPU calculation.

## Arc Length Calculation

The first algorithm, I decided to calculate the arc length that we used in our web front-end application. In our application, the calculation of arc length was completed by the D3 library, but the basic idea is simple: get the summation all numbers of reviews up, and divided each business's reviews with the sum, then multiply with the length of the whole circle.

So I wrote the Python code as follow and run in my EC2 instance. In this code, I just use a little part of original dataset since the capacity of GPU is limited.

```
import pycuda.gpuarray as gpuarray
import pycuda.driver as cuda
import pycuda.tools, pycuda.autoinit, pycuda.compiler
import numpy, math

start = cuda.Event()
end = cuda.Event()
d = [2.0, 3.0, 3.0, 3.0, 13.0, 5.0, 20.0, 5.0, 19.0, 12.0, 10.0, 17.0, 6.0, 2.0,
3.0, 2.0, 39.0, 14.0, 2.0, 31.0, 113.0, 33.0, 74.0, 61.0, 23.0, 21.0, 12.0, 4.0,
300.0, 8.0, 2.0, 89.0, 25.0, 2.0, 39.0, 7.0, 9.0, 2.0, 52.0, 39.0, 3.0, 4.0,
78.0, 2.0, 3.0, 5.0, 10.0, 10.0, 9.0, 4.0, 4.0, 19.0, 3.0, 2.0, 84.0, 43.0,
44.0, 2.0, 2.0, 3.0, 5.0, 25.0, 2.0, 16.0, 34.0, 4.0, 13.0, 68.0, 3.0, 2.0, 2.0,
25.0, 77.0, 5.0, 15.0, 5.0, 4.0, 4.0, 8.0, 10.0, 7.0, 8.0, 46.0, 27.0, 2.0, 8.0,
15.0, 10.0, 2.0, 61.0, 16.0, 99.0, 2.0, 13.0, 27.0, 36.0, 2.0, 5.0, 5.0, 8.0]
r = 500
a_gpu = gpuarray.to_gpu(numpy.array(d))

# Run on CPU
start.record()
r_cpu = 2 * math.pi * r / sum(d)
c_cpu = numpy.array([r_cpu * x for x in d])
end.record()
end.synchronize()

print "-" * 80
print "Coefficient of CPU operation:",
```

```

print r_cpu
print "CPU time: %fs" %(start.time_till(end) * 1e-3)
print c_cpu

# Run on GPU
start.record()
r_gpu = 2 * math.pi * r / gpuarray.sum(a_gpu).get()
c_gpu = (r_gpu * a_gpu).get()
end.record()
end.synchronize()

print "-" * 80
print "Coefficient of GPU operation:",
print r_gpu
print "GPU time: %fs" %(start.time_till(end) * 1e-3)
print c_gpu

print "-" * 80
print "CPU-GPU difference:"
print c_cpu - c_gpu

numpy.allclose(c_cpu, c_gpu)

```

The result is as follow. It's very interesting that the GPU is very much slower than CPU in this calculation. Perhaps the large scale of array effect the speed of GPU when it take a lot of time for copying data for host(CPU) to device(GPU), and, in the other hand, maybe this calculation couldn't take any advantage from GPU's ability.

But anyway, this two calculations get the same result.

```
ubuntu@ip-172-31-61-7:~/HW3$ python Example_1.py
```

```
-----
Coefficient of CPU operation: 1.43714211052
```

```
CPU time: 0.000035s
```

```
[  2.87428422  4.31142633  4.31142633  4.31142633 18.68284744
   7.18571055 28.74284221  7.18571055 27.3057001 17.24570533
  14.37142111 24.43141588  8.62285266  2.87428422  4.31142633
   2.87428422 56.04854231 20.11998955  2.87428422 44.55140543
 162.39705849 47.42568965 106.34851618 87.66566874 33.05426854
  30.17998432 17.24570533  5.74856844 431.14263316 11.49713688
   2.87428422 127.90564784 35.92855276  2.87428422 56.04854231
 10.05999477 12.93427899  2.87428422 74.73138975 56.04854231
   4.31142633  5.74856844 112.09708462  2.87428422  4.31142633
   7.18571055 14.37142111 14.37142111 12.93427899  5.74856844
   5.74856844 27.3057001  4.31142633  2.87428422 120.71993728
  61.79711075 63.23425286  2.87428422  2.87428422  4.31142633]
```

7.18571055	35.92855276	2.87428422	22.99427377	48.86283176
5.74856844	18.68284744	97.72566352	4.31142633	2.87428422
2.87428422	35.92855276	110.65994251	7.18571055	21.55713166
7.18571055	5.74856844	5.74856844	11.49713688	14.37142111
10.05999477	11.49713688	66.10853708	38.80283698	2.87428422
11.49713688	21.55713166	14.37142111	2.87428422	87.66566874
22.99427377	142.27706894	2.87428422	18.68284744	38.80283698
51.73711598	2.87428422	7.18571055	7.18571055	11.49713688]

-----

Coefficient of GPU operation: 1.43714211052

GPU time: 0.531286s

[	2.87428422	4.31142633	4.31142633	4.31142633	18.68284744
	7.18571055	28.74284221	7.18571055	27.3057001	17.24570533
	14.37142111	24.43141588	8.62285266	2.87428422	4.31142633
	2.87428422	56.04854231	20.11998955	2.87428422	44.55140543
	162.39705849	47.42568965	106.34851618	87.66566874	33.05426854
	30.17998432	17.24570533	5.74856844	431.14263316	11.49713688
	2.87428422	127.90564784	35.92855276	2.87428422	56.04854231
	10.05999477	12.93427899	2.87428422	74.73138975	56.04854231
	4.31142633	5.74856844	112.09708462	2.87428422	4.31142633
	7.18571055	14.37142111	14.37142111	12.93427899	5.74856844
	5.74856844	27.3057001	4.31142633	2.87428422	120.71993728
	61.79711075	63.23425286	2.87428422	2.87428422	4.31142633
	7.18571055	35.92855276	2.87428422	22.99427377	48.86283176
	5.74856844	18.68284744	97.72566352	4.31142633	2.87428422
	2.87428422	35.92855276	110.65994251	7.18571055	21.55713166
	7.18571055	5.74856844	5.74856844	11.49713688	14.37142111
	10.05999477	11.49713688	66.10853708	38.80283698	2.87428422
	11.49713688	21.55713166	14.37142111	2.87428422	87.66566874
	22.99427377	142.27706894	2.87428422	18.68284744	38.80283698
	51.73711598	2.87428422	7.18571055	7.18571055	11.49713688]

-----

CPU-GPU difference:

[	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.
	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.	0.]

ubuntu@ip-172-31-61-7:~/HW3\$

## Fibonacci Sequence

Because the first example cannot show the benefits of using GPU programming, this time I decided to implement a very famous problem in algorithm: the n-th Fibonacci sequence. Since its hard to call recursion in GPU, I'll use matrix approach to calculate it. The matrix approach's

formula is as follow:

$$\begin{bmatrix} F(n+1) \\ F(n) \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}^n \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

For the convenient of programming, I make all matrixes at 2x2 scale, and calculate the 30-th Fibonacci sequence in both CPU and GPU.

```
import numpy as np
import pycuda.autoinit
import pycuda.driver as cuda
from pycuda import compiler, gpuarray, tools

start = cuda.Event()
end = cuda.Event()
MATRIX_SIZE = 2
seq = 30
kernel_code_template = """
__global__ void MatrixMulKernel(float *a, float *b, float *c)
{
    // 2D Thread ID (assuming that only *one* block will be executed)
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue is used to store the element of the matrix
    // that is computed by the thread
    float Pvalue = 0;

    // Each thread loads one row of M and one column of N,
    // to produce one element of P.
    for (int k = 0; k < %(MATRIX_SIZE)s; ++k) {
        float Aelement = a[ty * %(MATRIX_SIZE)s + k];
        float Belement = b[k * %(MATRIX_SIZE)s + tx];
        Pvalue += Aelement * Belement;
    }

    // Write the matrix to device memory;
    // each thread writes one element
    c[ty * %(MATRIX_SIZE)s + tx] = Pvalue;
}
"""

# Run on CPU
a_cpu = np.matrix('1 1; 1 0').astype(np.float32)
b_cpu = np.matrix('1 0; 1 0').astype(np.float32)
c_cpu = a_cpu

start.record()
```

```

for x in range(0, 29):
    c_cpu = c_cpu * a_cpu
c_cpu = c_cpu * b_cpu
end.record()
end.synchronize()

print "-" * 80
print "CPU time: %fs" %(start.time_till(end) * 1e-3)
print "Fibonacci(30) by CPU:",
print c_cpu.item(2)

# Run on GPU
a_gpu = gpuarray.to_gpu(a_cpu)
b_gpu = gpuarray.to_gpu(b_cpu)
c_gpu = gpuarray.to_gpu(a_cpu)

kernel_code = kernel_code_template % { 'MATRIX_SIZE': MATRIX_SIZE }
mod = compiler.SourceModule(kernel_code)
matrixmul = mod.get_function("MatrixMulKernel")

start.record()
for x in range(0, 29):
    matrixmul(c_gpu, a_gpu, c_gpu, block = (MATRIX_SIZE, MATRIX_SIZE, 1),)
matrixmul(c_gpu, b_gpu, c_gpu, block = (MATRIX_SIZE, MATRIX_SIZE, 1),)
end.record()
end.synchronize()

print "-" * 80
print "GPU time: %fs" %(start.time_till(end) * 1e-3)
print "Fibonacci(30) by GPU:",
print c_gpu.get().item(2)

print "-" * 80
print "CPU-GPU difference:"
print c_cpu.item(2) - c_gpu.get().item(2)

np.allclose(c_cpu, c_gpu.get())

```

Although the performance of GPU is still not better than CPU, the time difference between them is smaller.

```
ubuntu@ip-172-31-61-7:~/HW3$ python Example_2.py
```

```
-----  
CPU time: 0.000251s
```

```
Fibonacci(30) by CPU: 1346269.0
```

```
-----  
GPU time: 0.001506s
```

```
Fibonacci(30) by GPU: 1346269.0
```

```
-----  
CPU-GPU difference:
```

```
0.0
```

```
ubuntu@ip-172-31-61-7:~/HW3$
```